


April 30, 2013




IT trends


- ▶ More resources
 - ▶ Better hardware at lower costs
 - ▶ Higher standards for software quality
- ▶ More users
 - ▶ Contact with technology at an earlier age
 - ▶ Shared access to the same device
- ▶ Data consolidation
 - ▶ Data warehousing
 - ▶ Service unification
 - ▶ Differentiated access
- ▶ Increased flexibility
 - ▶ Versatile configuration
 - ▶ Focus on usability

Kernel Support

- ▶ Namespaces:
 - ▶ Abstract resources
 - ▶ Processes see the resource as their own
 - ▶ Isolation between namespaces



- ▶ Control Groups
 - ▶ Resource management among processes
 - ▶ Hierarchical support
 - ▶ Interaction with resource responsible structures:
 - ▶ Scheduler
 - ▶ Pager



Sample Process Hierarchy

```
init(1)-+--dnsmasq(2162)
|
|-klogd(2175)
|-lxc-start(2964)---init(2966)-----+--init(2972)
|                                     |-sh(2971)
|                                     '-syslogd(2969)
|
|
|
|-lxc-start(2974)---init(2976)-----+--init(2982)
|                                     |-sh(2981)
|                                     '-syslogd(2979)
|
|
|-netserver(2167)
|-sh(2179)
|-syslogd(2173)
'-udevd(962)-+--udevd(1189)
              '-udevd(1190)
```



FreeScale, the FreeScale logo, i.MX, eC, C-5, Code TEST, CodeSimulator, ColdFire, C-4196, the Energy Efficient Solutions logo, mbedG T, PowerQUICC, QorD, iZorC and Symphony are trademarks of Freescale Semiconductor, Inc., Reg. U.S. Pat. & Tm. O.F. Bell, BeeStar, ColdFire, CoreJet, Rebs, iNels, MFC, PABlockin a Package, Processor Expert, QorD Convergence, Cortex, QVCC Engine, SMM T.D.S, TubeLink, JorD and iZorC are trademarks of Freescale Semiconductor, Inc. All other product names are the property of their respective owners. © 2008 Freescale Semiconductor, Inc.

Process IDs

```
init(1) -- dnsmasq(2162)
| -klogd(2175)
| -lxc-start(2964) --- init(2966) (1) -- init(2972) (7)
|                                     | -sh(2971) (6)
|                                     ' -syslogd(2969) (4)
|
|
| -lxc-start(2974) --- init(2976) (1) -- init(2982) (7)
|                                     | -sh(2981) (6)
|                                     ' -syslogd(2979) (4)
|
|
| -netserver(2167)
| -sh(2179)
| -syslogd(2173)
' -udevd(962) -- udevd(1189)
      ' -udevd(1190)
```



Namespace Segregation

```
init(1)--dnsmasq(2162)
|-klogd(2175)
|-lxc-start(2964)--- init(2966)(1)--init(2972)(7)
|                                     |-sh(2971)(6)
|                                     '-syslogd(2969)(4)
|                                     PID Namespace 1
|
|-lxc-start(2974)--- init(2976)(1)--init(2982)(7)
|                                     |-sh(2981)(6)
|                                     '-syslogd(2979)(4)
|                                     PID Namespace 2
|
|-netserver(2167)
|-sh(2179)
|-syslogd(2173)
'-udevd(962)--udevd(1189)
                '-udevd(1190)
```



FreeScale, the FreeScale logo, 32/64, C-4, CodeWISIT, CodeLinker, ColdFire, C-Move, the Energy Efficient Solutions logo, iM8000 T, PowerQUICC, QorIQ, iZorQ and Symphony are trademarks of Freescale Semiconductor, Inc., Reg. U.S. Pat. & Tm. Off. BeeR, BeeStad, ColdFire, CoolNet, Hots, iZorQ, J4C, P4000 are a Package, Processor Expert, QorIQ Converge, QorIQ, QorIQ Engine, QorIQ T400, TuboLink, iZorQ and iZorQ are trademarks of Freescale Semiconductor, Inc. All other product or service names are the property of their respective owners. © 2014 Freescale Semiconductor, Inc.

"chroot on steroids"

```
root: /var/lib/lxc/foo1/rootfs/
```

```
init(2966)(1) -+- init(2972)(7)
```

$$| -sh(2971) (6)$$

```
'-syslogd(2969) (4)
```

PID Namespace 1

```
root: /var/lib/lxc/foo1/rootfs/
```

```
init(2976)(1) -+- init(2982)(7)
```

$$| -sh(2981) (6)$$

```
'-syslogd(2979) (4)
```

PID Namespace 2

CPU Partitioning

```

init(1)--dnsmasq(2162)
    |-klogd(2175)
,-----|-lxc-start(2964)---
|         |                25%
|         |
|         |
|         |
|         |
1 core    |-lxc-start(2974)---
|         |                75%
|         |
|         |
|         |
'-----|-----
        |-netserver(2167)
        |-sh(2179)
        |-syslogd(2173)
        '-udevd(962)--udevd(1189)
                '-udevd(1190)

```

Demo

1. Start 2 containers
2. Check PIDs
3. Assign them a single core on the host
4. Balance CPU usage 25% - 75%




FreeScale, the FreeScale logo, ARM, Cx, CodeWIT, CodeLinker, ColdFire, C-More, the Energy Efficient Solutions logo, iM8000 T, PowerQUICC, QorIQ, iZaQ and Symphony are trademarks of FreeScale Semiconductor, Inc., Reg. U.S. Pat. & Tm. Off. Bell, Be!lStad, ColdFire+, CoolNet, Flex, iFlex, iMCC, P4tterns a Package, Processor Expert, QorIQ Converge, QorIQ, QVCC Engine, S3AR T4D5, TubuLink, iVfora and iVtronic are trademarks of FreeScale Semiconductor, Inc. All other product or service names are the property of their respective owners. © 2011 Freescale Semiconductor, Inc.


System Performance

CPU performance


Linpack



Mem throughput




Bandwidth NetPIPE



Latency

NetPIPE



Isolation

PERFORMANCE ISOLATION FOR LU APPLICATION. THE RESULTS REPRESENT HOW MUCH THE APPLICATION PERFORMANCE IS IMPACTED BY DIFFERENT STRESS TESTS IN ANOTHER VM/CONTAINER. DNR MEANS THAT APPLICATION WAS NOT ABLE TO RUN.

	LXC	OpenVZ	VServer	Xen
CPU Stress	0	0	0	0
Memory	88.2%	89.3%	20.6%	0.9%
Disk Stress	9%	39%	48.8%	0
Fork Bomb	DNR	0	0	0
Network Receiver	2.2%	4.5%	13.6%	0.9%
Network Sender	10.3%	35.4%	8.2%	0.3%

Thank you!

Questions?



FreeScale, the FreeScale logo, i.MX, C-5, CodeWISIT, CodeLinker, ColdFire, Co-More, the Energy Efficient Solutions logo, iM800 T, PowerQUICC, QorD, iLibra and Symphony are trademarks of Freescale Semiconductor, Inc., Reg. U.S.Pat. & Tm.O.F. Bell E, BeeStar, ColdFire+, CoolJet, Flex, I, Inels, JEEC, PABRish A Package, Processor Expert, QorD Converge, Onviva, QUICC Engine, S&S T.D.S, TubeLink, Viorio and iLibric are trademarks of Freescale Semiconductor, Inc. All other product or service names are the property of their respective owners. © 2011 Freescale Semiconductor.